

Introduction to R for Stata Users

01: Introduction

Purushottam Mohanty

03 June 2021

Why switch to R?

While this is not an exhaustive list, the primary reasons can be listed as:

- R is a free and open source software making it accessible for students and researchers across the world. (goodbye to license worries!)
- R is way faster and resource efficient.
- Syntax much closer to other general programming languages such as Python and can be efficiently scaled up.
- Allows users to take advantage of thousands of packages developed and maintained by a large base of users.
- R is widely used outside academia, thus opening up industry based opportunities in data science and related fields.
- Strong community of developers and enthusiasts across multiple platforms. (quick stackoverflow solutions!)
- R Markdown allows writing and exporting latex documents, beamer presentations and publish them as HTML or PDFs within R.

Installing R and RStudio

The latest version of **R** can be installed from The Comprehensive R Archive Network (CRAN).

1. Download R for [macOS](#).¹
2. Download R for [Windows](#).
3. Download R for [Linux](#).

In order to take complete advantage of R, download **RStudio** too. RStudio is a GUI for R that provides code editor, data and graphics viewer, git integration and command line access along with various other features.

1. Download [RStudio](#).²

¹ **Note:** Download the Intel installer even if you have a M1 Mac as many libraries haven't been updated to use the ARM version.

² **Note:** Download the RStudio Desktop free version. Available for macOS, Windows and Linux.

Basics of R

1. R is an object oriented language. Every dataset, figure, text, string, number or regression result is or can be stored as an object.
2. Depending on what is stored in an object, the `class` of an object is defined which correspondingly determines the operations can be made on the object. (eg: cannot calculate mean of a string)
3. Each object is given a namespace using which we store or call an object. (eg: `x = 1` where `x` is an object.) Namespaces have to be unique or they are overwritten.
4. All operations are performed using functions which can be user-written or accessed from a library.
5. R has different libraries (i.e. packages) which enable additional functions and features beyond the `baseR` library.

Key Differences from Stata

1. In Stata `ssc install` is run once to install and access a particular package whereas in R packages can to be installed once but loaded into memory everytime.
2. In R one can use multiple dataframes (i.e. datasets) simultaneously so no need to constantly dabble with `preserve`, `restore` and `temp save`.
3. Since, everything is stored as objects in R one easily access elements within objects through consistent syntax.
4. While `putdocx` in Stata can be helpful, RMarkdown provides much easier ways to write latex documents or beamer presentations and export them as PDFs or html.
5. R and RStudio are free and updated regularly so no need to be stuck with older versions (hello Stata 14 users!).

Additional Resources

1. [R for Data Science](#) - by Hadley Wickhman and Garrett Grolemond.
2. [R Graphics: Codebook](#) - by Winston Chang
3. [Introduction to Data Science](#) - by Rafael A. Irizarry
4. [Data Science for Economists](#) - by Grant McDermott
5. [GSU Library Course](#)
6. [R Markdown: The Definitive Guide](#) - by Yihui Xie, J. J. Allaire, Garrett Grolemond

Installing Packages

First, open RStudio > New File > R Script and save it in the directory of choice. Let's install the `tidyverse` and `gapminder` packages using,

```
install.packages('tidyverse') # composite of data cleaning and visualization packages  
install.packages('gapminder') # socio-economic data for countries
```

After installation, the packages need to be loaded into memory. Note that packages which have been installed earlier need not be installed again and can directly be loaded into memory.

```
library(tidyverse)  
library(gapminder)
```

A more effective way to do the installation and loading process is using the `pacman` package. Just install the package once using `install.packages('pacman')` and then use,

```
pacman::p_load(tidyverse, gapminder)
```

The `PACKAGE::function()` syntax allows one to use functions from installed packages without loading them.

Object Types

R objects can be of different `class` depending on the type of data that is assigned to them.

```
# character
x = "London"
# numeric
x = 100
# list
x = as.list(c("London", "Paris", "Zurich", "Madrid", "Amsterdam"))
# dataframe
x = read_csv("./example.csv")
```

There are a large number of object types used in R but `character`, `numeric`, `list`, `dataframe` and `matrix` are the most commonly used ones. Objects can also be transformed between types.

The object type determines the operations that can be performed on the object.

```
x = 1
y = "A"
x + y # cannot be performed (numeric + character)
```


Reading Files

Before we proceed to reading different file formats, let's set the base directory first,

```
getwd() # getting current directory
setwd('/Users/purushottam/Documents/GitHub/prog-notes/stataR') # setting directory

# comma delimited file
library(readr)
read_delim(file = './example.csv', delim = ',')
# excel file (reads .xls and .xlsx files)
library(readxl)
read_excel(path = './example.xlsx', sheet = 'sheet1')
# stata file
library(haven)
read_stata(file = './example.dta')
# json file
library(jsonlite)
read_json(path = './example.json')
```

Similarly, there are other packages to load SAS, SPSS, TIFF or other files.

Dataframes

Datasets in R are called `dataframe`. If you tried any of the functions listed earlier to read files you would notice that while they appear in the console, there is no way to access them. In order to do so, we need to store them as an object. This can be done using,

```
df ← read_stata(file = './example.dta')
```

The `.dta` file is now stored as a `dataframe` class object `df` and can be viewed either by selecting it in the environment window in R Studio or by simply executing `df`. Note that many R users like to use `=` instead of `←` as the assign function. They both work the same so pick what you want but stay consistent.

Datasets can either be imported into R as `dataframe` or it can be created in the following way.

```
df = data.frame(names = c("London", "Paris", "Zurich", "Madrid", "Amsterdam"),  
               rents = c(1800, 1400, 2000, 1000, 1700))
```

Dataframes (cont.)

For the current tutorial however we shall data from `gapminder` package. The package provides life expectancy, GDP per capita, and population data at the country level for every 5 years from 1952 to 2007.

```
df = gapminder
df
```

```
## # A tibble: 1,704 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

Indexing

All elements within objects in R can be accessed using row and column index or names. R follows a consistent form of indexing which is applicable to all class of objects.

A `dataframe` in R can be indexed by using the following syntax `dataframeName[rowNumber, columnNumber]` or `dataframeName[rowNumber, columnName]`.

```
df[1,4]
```

```
## # A tibble: 1 x 1
##   lifeExp
##   <dbl>
## 1    28.8
```

```
df[1, 'lifeExp']
```

```
## # A tibble: 1 x 1
##   lifeExp
##   <dbl>
## 1    28.8
```

Indexing (cont.)

Multiple rows or columns can also be indexed.

```
df[1, 4:5] # row 1 and column 4-5
```

```
## # A tibble: 1 x 2
##   lifeExp    pop
##   <dbl>   <int>
## 1   28.8 8425333
```

```
df[1, c('lifeExp', 'pop')]
```

```
## # A tibble: 1 x 2
##   lifeExp    pop
##   <dbl>   <int>
## 1   28.8 8425333
```

Indexing can also be done for only rows or columns.

```
df[1,] # row 1 - all columns
df[,4] # all rows - column 4
df[4] # all rows - column 4
```